

Analysis of car accidents in Mexico City

2020-03-22

1 Intro

```
rm(list = ls())  
#Load packages  
library(dplyr) #v. 0.8.5  
library(htmlwidgets) #v. 1.5.1  
library(mgcv) #v. 1.8-28  
library(plotmo) #v. 3.5.6  
library(randomForest) #v. 4.6-14  
library(ranger) #v. 0.12.1  
library(RColorBrewer) #v. 1.1-2  
# devtools::install_github("hrbrmstr/streamgraph")  
library(streamgraph) #v. 0.9.0  
library(xtable) #v. 1.8-4
```

Some colors to start with

```
COL <- c("black",  
        rgb(100, 38, 33, maxColorValue = 100), #red  
        rgb(0, 65, 55, maxColorValue = 100), #green  
        rgb(28, 24, 61, maxColorValue = 100), #blue  
        rgb(76, 32, 72, maxColorValue = 100), #purple  
        rgb(21, 75, 87, maxColorValue = 100), #cyan  
        rgb(0, 47, 59, maxColorValue = 100) #dark cyan  
)
```

2 Data description

```
#load data  
load("./dataderived/image_preprocessBoth.RData")  
  
#load accidents data again to see the types of accidents  
DA <- read.csv("./dataraw/accidents.csv", nrow = 155466)  
names(DA)[c(2, 3, 7, 11)] <- c("Month", "Year", "Day", "Type")  
#Truncate to the period of analysis (whole years here):  
DA <- DA[(DA$Year >= 2001) & (DA$Year <= 2015),]  
#DA <- DA[!(DA$Year == 2015 & DA$Month == 12),] #if need to remove Dec 2015
```

Number of accidents by year and type

The codes for the types of accidents:

1. Collision with other vehicle
2. Collision with pedestrian

3. Collision with animal
4. Collision with fixed object
5. Flip
6. Passenger fall off
7. Drive to ditch
8. Fire
9. Collision with train
10. Collision with motorcycle
11. Collision with bicycle
12. Other

Count accidents by year and type

```
DA$Count <- 1L
da <- aggregate(DA$Count, by = list(DA$Year, DA$Type), FUN = sum)
names(da) <- c("Year", "Type", "Count")
```

Number of different types of accidents

```
length(unique(da$Type))

## [1] 12
```

Figure 2: Number of car accidents per year in Mexico City

```
tmp <- brewer.pal(11, name = "Spectral")
#Add 1 more color to this palette of 11 colors and rearrange for a better look
tmp[5] <- "black"
tmp[7] <- tmp[1]
COL2 <- c(tmp[-1], COL[c(2, 5)])
pp = streamgraph(da, "Type", "Count", "Year",
  offset = "zero", order = "asis",
  interactive = TRUE) %>%
  sg_axis_x(1, "Year", "%Y") %>%
  sg_fill_manual(COL2) # sg_fill_brewer("Spectral")
#Save the widget then print in PDF and edit for the paper
saveWidget(pp, file = "AccidTypes.html")
```

Percentage by type

```
typeA <- table(DA$Type)
round(typeA * 100 / sum(typeA), 1)

##
##      1      2      3      4      5      6      7      8      9     10     11     12
## 65.6  5.5  0.3  8.9  4.8  0.5  4.0  0.1  0.0  8.0  1.3  1.0
```

Number by year

```
tmp = table(DA$Year)
tmp

##
## 2001 2002 2003 2004 2005 2006 2007 2008 2009 2010 2011 2012 2013 2014 2015
## 7473 8024 8057 8858 8731 9396 7611 10214 10197 8711 8523 9431 10814 9502 9294
```

Average (percentage) increase by year

```
mean(diff(tmp)) #average increase

## [1] 130
```

```

MeanRelChange = (tmp[length(tmp)] / tmp[1]) ^ (1 / (length(tmp) - 1) )
MeanRelChange*100 - 100 #average percentage increase

## 2015
## 1.57

#check should be close to 0:
tmp[1] * MeanRelChange^((length(tmp) - 1)) - tmp[length(tmp)]

##      2001
## -1.09e-11

```

Collisions with motorcycles

```

tmp = da$Count[da$Type == 10] #select by type
tmp = tmp[c(1, length(tmp))] #select 1st and last years
tmp

## [1] 378 1320

tmp[2]/tmp[1] #increase times

## [1] 3.49

```

Collisions with pedestrians

```

tmp = da$Count[da$Type == 2] #select by type
tmp = tmp[c(1, length(tmp))] #select 1st and last years
tmp

## [1] 1455 377

tmp[1]/tmp[2] #decrease times

## [1] 3.86

```


Number of cars registered

Figure 3: Number of cars registered in Mexico City, 2001–2015

```

CR <- CR[CR$Year >= 2001 & CR$Year <= 2015, ]
for(i in 1:2){
  if(i == 1) pdf("./figures/tsCarsReg.pdf", width = 8, height = 3.7)
  par(mar = c(3.5, 3.5, 0.5, 0.1), mgp = c(2.5, 1, 0))
  plot(CR$Year, CR$CarsReg/1000000,
       xlab = "Year", ylab = "Number of cars, million",
       ylim = c(0, 5), yaxs="i",
       type = "o", pch = 16, cex = 1.5,
       panel.first = grid(nx = 0, ny = 5, lty = 1),
       col = COL[1], lty = 1, las = 1, lwd = 2)
  if(i == 1) dev.off()
}

```



Average (percentage) increase by year

```
tmp = CR$CarsReg/1000000 #cars registered, million
tmp[c(1, length(tmp))] #select 1st and last years

## [1] 2.23 4.63

mean(diff(tmp)) #average increase

## [1] 0.171

MeanRelChange = (tmp[length(tmp)] / tmp[1]) ^ (1 / (length(tmp) - 1))
MeanRelChange*100 - 100 #average percentage increase

## [1] 5.34

#check should be close to 0:
tmp[1] * MeanRelChange^((length(tmp) - 1)) - tmp[length(tmp)]

## [1] 8.88e-16
```

Percent missing values


```
tmp = apply(is.na(DataHour), 2, mean)
max(tmp * 100)

## [1] 0.0635
```


Weather

Figure 4: Time series plots of daily accident rate, total rainfall, and average air temperature


```
D <- DataDay
for(i in 1:2){
  if(i == 1) pdf("./figures/tsNaccid.pdf", width = 8, height = 3.7)
  par(mar = c(3.5, 3.5, 0.1, 0.1), mgp = c(2.5, 1, 0))
  plot(D$Date, D$NAccidPer100000,
       xlab = "Date",
       ylab = "Number of accidents per 100,000 cars",
       col = COL[2], type = "l", lty = 1, las = 1)
  if(i == 1) dev.off()
}
```



```
for(i in 1:2){
  if(i == 1) pdf("./figures/tsRain.pdf", width = 8, height = 3.7)
  par(mar = c(3.5, 3.5, 0.1, 0.1), mgp = c(2.5, 1, 0))
  plot(D$Date, D$Rain,
       xlab = "Date",
       ylab = "Rainfall per day (mm)",
       col = COL[4], type = "l", lty = 1, las = 1)
  if(i == 1) dev.off()
}
```




```
for(i in 1:2){
  if(i == 1) pdf("./figures/tsTemp.pdf", width = 8, height = 3.7)
  par(mar = c(3.5, 3.5, 0.1, 0.1), mgp = c(2.5, 1, 0))
  plot(D$Date, D$Temperature,
       xlab = "Date",
       ylab = expression(paste("Average temperature per day (", degree, "C)")),
       col = COL[5], type = "l", lty = 1, las = 1)
  if(i == 1) dev.off()
}
```




Weather – for conclusions

Weather by hour

```
tmp <- tapply(DataHour$Temperature, DataHour$Hour, mean, na.rm = TRUE)
plot(tmp, x = 0:23, type = "o", panel.first = grid(lty = 1),
      ylab = "Average temperature per hour",
      las = 1)
#add lines for 'busiest' hours
abline(v = c(14.5, 16.5), col = 2, lty = 2)
tmp <- tapply(DataHour$Rain, DataHour$Hour, mean, na.rm = TRUE)
plot(tmp, x = 0:23, type = "o", panel.first = grid(lty = 1),
      ylab = "Average rainfall per hour",
      las = 1)
#most rainy hours
abline(v = c(18, 21), col = 2, lty = 2)
```





Rainy months (apply it to full years only, i.e., before 2015)

```
tapply(DataDay$Rain[DataDay$Year < 2015], DataDay$Month[DataDay$Year < 2015], mean) *
length(unique(DataDay$Year[DataDay$Year < 2015]))
```

##	1	2	3	4	5	6	7	8	9	10	11	12
##	1.086	1.915	2.318	6.377	10.359	30.877	25.009	27.180	26.418	13.738	4.500	0.587

3 Methods

4 Results

Table 1: Quartile summaries (daily data)

The quartiles table with std dev. for the mean (divide by \sqrt{n})

```
D <- DataDay
labs = seq(0.25, 1, by = 0.25)
#temperature quartiles
tq = quantile(D$Temperature, probs = c(0, 0.25, 0.50, 0.75, 1))
tq

##      0%    25%    50%    75%   100%
##  6.31 15.41 16.87 18.18 23.85

D$tempInt = cut(D$Temperature, breaks = tq, labels = paste("temp_", labs, sep = ""),
               include.lowest = TRUE, right = FALSE, ordered_result = TRUE)
#rain quartiles (FOR RAINY DAYS!)
rq = quantile(D$Rain[D$Rain > 0], probs = c(0, 0.25, 0.50, 0.75, 1))
rq[1] = 0
rq

##      0%    25%    50%    75%   100%
##  0.000  0.178  0.810  2.200 76.870

D$rainInt = cut(D$Rain, breaks = rq, labels = paste("rain_", labs, sep = ""),
               include.lowest = TRUE, right = FALSE, ordered_result = TRUE)
#Summary per intersection of the quartiles:
magg1 <- tapply(D$NAccidPer100000, list(D$tempInt, D$rainInt), mean)
magg1 <- format(round(magg1, 2), digits = 2)
#Sample size per intersection of the quartiles:
ss <- table(D$tempInt, D$rainInt)
sdagg1 <- tapply(D$NAccidPer100000, list(D$tempInt, D$rainInt), sd)
sdagg1 <- sdagg1 / sqrt(ss)
```

```

sdagg1 <- format(round(sdagg1, 2), digits = 2)
M <- matrix(paste(magg1, "(", sdagg1, ")", sep = ""), nrow = 4)
dimnames(M) <- dimnames(magg1)
#Copy this from R console into latex:
print(xtable(M,
             caption = "Average accident rate, st.dev. in the parentheses",
             label = "tab:TempRain", size = "small"))

## % latex table generated in R 3.6.1 by xtable 1.8-4 package
## % Sun Mar 22 03:52:24 2020
## \begin{table}[ht]
## \centering
## \begin{tabular}{rllll}
## \hline
## & rain\_0.25 & rain\_0.5 & rain\_0.75 & rain\_1 \\
## \hline
## temp\_0.25 & 0.79 (0.01) & 0.80 (0.03) & 0.85 (0.04) & 0.80 (0.03) \\
## temp\_0.5 & 0.78 (0.01) & 0.79 (0.02) & 0.80 (0.02) & 0.79 (0.02) \\
## temp\_0.75 & 0.75 (0.01) & 0.74 (0.02) & 0.74 (0.02) & 0.79 (0.02) \\
## temp\_1 & 0.74 (0.01) & 0.69 (0.02) & 0.72 (0.02) & 0.73 (0.03) \\
## \hline
## \end{tabular}
## \caption{Average accident rate, st.dev. in the parentheses}
## \label{tab:TempRain}
## \end{table}

```

4.1 GAM

Daily GAM

```

D <- DataDay
D$Month <- factor(D$Month)
# Create train+test data
DtrainDay <- D[D$Year <= 2012,]
DtestDay <- D[D$Year > 2012,]
# summary(DtrainDay)
# summary(DtestDay)

```

Size of the training and testing data

```

nrow(DtrainDay)

## [1] 4383

nrow(DtestDay)

## [1] 1064

K <- 5
set.seed(140)
gamfit <- gamDay <- mgcv::gam(NAccidPer100000 ~ s(Year, k = K)
+ Month
+ Weekday
+ HNSSaturday + Holiday
+ te(Rain, Temperature, k = K)
, select = TRUE
, bs = "cr"
, method = "REML"
, data = DtrainDay)

anova(gamfit)

##
## Family: gaussian
## Link function: identity
##
## Formula:

```



```

## NAccidPer100000 ~ s(Year, k = K) + Month + Weekday + HNSSaturday +
##      Holiday + te(Rain, Temperature, k = K)
##
## Parametric Terms:
##           df      F p-value
## Month      11  8.08 3.9e-14
## Weekday     6 93.95 < 2e-16
## HNSSaturday 1  6.27 0.01229
## Holiday     1 12.35 0.00045
##
## Approximate significance of smooth terms:
##           edf Ref.df      F p-value
## s(Year)      3.89   4.00 667.07 < 2e-16
## te(Rain,Temperature) 4.25  24.00   1.06 3.6e-06

summary(gamfit)

##
## Family: gaussian
## Link function: identity
##
## Formula:
## NAccidPer100000 ~ s(Year, k = K) + Month + Weekday + HNSSaturday +
##      Holiday + te(Rain, Temperature, k = K)
##
## Parametric coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.770953   0.014288  53.96 < 2e-16 ***
## Month2       0.046945   0.015805   2.97 0.00299 **
## Month3       0.028194   0.016555   1.70 0.08863 .
## Month4      -0.000206   0.017905  -0.01 0.99083
## Month5       0.039739   0.018114   2.19 0.02830 *
## Month6       0.027002   0.017511   1.54 0.12314
## Month7      -0.066507   0.016828  -3.95 7.9e-05 ***
## Month8      -0.021090   0.017015  -1.24 0.21524
## Month9      -0.000971   0.016895  -0.06 0.95417
## Month10     -0.000934   0.015994  -0.06 0.95344
## Month11     0.027360   0.015187   1.80 0.07168 .
## Month12     0.032281   0.014902   2.17 0.03035 *
## Weekday2    -0.062866   0.011493  -5.47 4.8e-08 ***
## Weekday3    -0.031328   0.011502  -2.72 0.00648 **
## Weekday4    -0.021796   0.011486  -1.90 0.05782 .
## Weekday5     0.063075   0.011480   5.49 4.1e-08 ***
## Weekday6     0.180532   0.013327  13.55 < 2e-16 ***
## Weekday7     0.112156   0.011473   9.78 < 2e-16 ***
## HNSSaturday1 -0.044911   0.017931  -2.50 0.01229 *
## Holiday1    -0.058966   0.016781  -3.51 0.00045 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##           edf Ref.df      F p-value
## s(Year)      3.89     4 667.07 < 2e-16 ***
## te(Rain,Temperature) 4.25  24   1.06 3.6e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) = 0.471  Deviance explained = 47.4%
## -REML = -695.15  Scale est. = 0.041158  n = 4383

concurvity(gamfit)

##           para s(Year) te(Rain,Temperature)
## worst      0.959   0.164               0.686
## observed   0.959   0.125               0.313
## estimate   0.959   0.120               0.060

# gam.check(gamfit) #commented out because plots are slow to render in PDF

```

```
# plot(gamfit)
# acf(residuals(gamfit, type = "pearson"), las = 1) #significant but low
```

Hourly GAM

```
D <- DataHour
D$Temperature.l1 <- dplyr::lag(D$Temperature, 1)
D$Rain.l1 <- dplyr::lag(D$Rain, 1)
D$Month <- factor(D$Month)
# Create train+test data
DtrainHour <- D[D$Year <= 2012,]
DtestHour <- D[D$Year > 2012,]
# summary(DtrainHour)
# summary(DtestHour)
```

Size of the training and testing data

```
nrow(DtrainHour)
```

```
## [1] 105192
```

```
nrow(DtestHour)
```

```
## [1] 25536
```

```
K <- 5
set.seed(140000)
gamfit <- gamHour <- mgcv::gam(NAccidPer100000 ~ s(Year, k = K)
+ Month + Weekday
+ HNSSaturday + Holiday
+ s(Hour, k = K)
+ te(Rain, Temperature, k = K)
+ te(Rain.l1, Temperature.l1, k = K)
, select = TRUE
, bs = "cr"
, method = "REML"
, data = DtrainHour)
```

```
anova(gamfit)
```

```
##
## Family: gaussian
## Link function: identity
##
## Formula:
## NAccidPer100000 ~ s(Year, k = K) + Month + Weekday + HNSSaturday +
##      Holiday + s(Hour, k = K) + te(Rain, Temperature, k = K) +
##      te(Rain.l1, Temperature.l1, k = K)
##
## Parametric Terms:
##              df      F p-value
## Month        11  12.32 < 2e-16
## Weekday       6 115.72 < 2e-16
## HNSSaturday   1   8.07  0.0045
## Holiday       1  15.21 9.6e-05
##
## Approximate significance of smooth terms:
##              edf Ref.df      F p-value
## s(Year)        3.91   4.00 841.9 <2e-16
## s(Hour)         3.99   4.00 851.8 <2e-16
## te(Rain,Temperature)  9.12  24.00  18.9 <2e-16
## te(Rain.l1,Temperature.l1) 6.09  24.00  13.9 <2e-16
```

```
summary(gamfit)
```

```
##
## Family: gaussian
## Link function: identity
```

```
##
## Formula:
## NAaccidPer100000 ~ s(Year, k = K) + Month + Weekday + HNSSaturday +
##      Holiday + s(Hour, k = K) + te(Rain, Temperature, k = K) +
##      te(Rain.l1, Temperature.l1, k = K)
##
## Parametric coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3.28e-02  5.12e-04  64.06 < 2e-16 ***
## Month2       1.58e-03  5.81e-04   2.71  0.0067 **
## Month3       6.48e-04  5.86e-04   1.11  0.2689
## Month4      -8.68e-04  6.17e-04  -1.41  0.1595
## Month5       5.97e-04  6.20e-04   0.96  0.3356
## Month6       7.33e-05  6.13e-04   0.12  0.9048
## Month7      -3.70e-03  5.95e-04  -6.21 5.2e-10 ***
## Month8      -1.86e-03  5.99e-04  -3.10  0.0019 **
## Month9      -9.72e-04  5.99e-04  -1.62  0.1043
## Month10     -7.75e-04  5.80e-04  -1.34  0.1812
## Month11      7.87e-04  5.67e-04   1.39  0.1652
## Month12      1.17e-03  5.59e-04   2.09  0.0364 *
## Weekday2    -2.58e-03  4.31e-04  -5.97 2.3e-09 ***
## Weekday3    -1.30e-03  4.32e-04  -3.01  0.0026 **
## Weekday4    -9.19e-04  4.31e-04  -2.13  0.0330 *
## Weekday5     2.62e-03  4.31e-04   6.08 1.2e-09 ***
## Weekday6     7.53e-03  5.00e-04  15.06 < 2e-16 ***
## Weekday7     4.70e-03  4.31e-04  10.91 < 2e-16 ***
## HNSSaturday1 -1.91e-03  6.73e-04  -2.84  0.0045 **
## Holiday1    -2.45e-03  6.29e-04  -3.90 9.6e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##              edf Ref.df      F p-value
## s(Year)         3.91     4 841.9 <2e-16 ***
## s(Hour)         3.99     4 851.8 <2e-16 ***
## te(Rain,Temperature)  9.12    24  18.9 <2e-16 ***
## te(Rain.l1,Temperature.l1) 6.09    24  13.9 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.124   Deviance explained = 12.4%
## -REML = -1.9638e+05   Scale est. = 0.0013898   n = 105095

concurvity(gamfit)

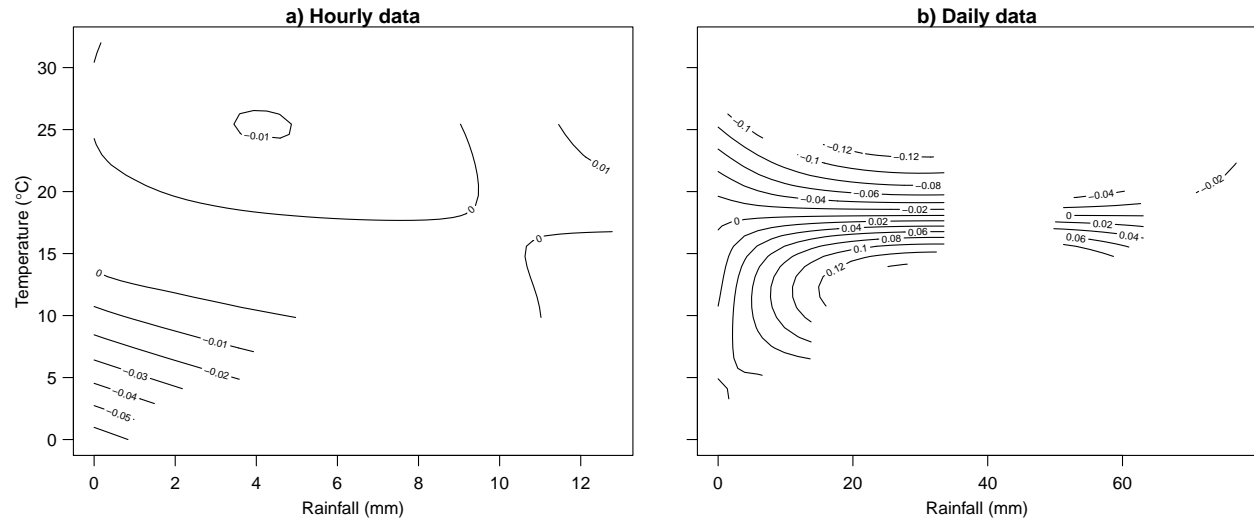
##              para s(Year) s(Hour) te(Rain,Temperature) te(Rain.l1,Temperature.l1)
## worst          0.95  0.144  0.709                0.965                0.964
## observed 0.95    0.109  0.660                0.890                0.908
## estimate 0.95    0.103  0.424                0.731                0.739

# gam.check(gamfit) #commented out because plots are slow to render in PDF
# plot(gamfit)
# acf(residuals(gamfit, type = "pearson"), las = 1) #significant but low
```

Figure 5: Contour plots of the tensor smooth terms

```
for(i in 1:2){
  if(i == 1) pdf("./figures/teTempRain.pdf", width = 8, height = 3.7)
  par(mar = c(3.3, 3, 1.1, 0.1), mgp = c(2.0, 0.8, 0))
  par(mfrow = c(1, 2))
  plot(gamHour, select = 3, se = FALSE, rug = FALSE, las = 1, ylim = c(0, 32), main = "a) Hourly data",
       xlab = "Rainfall (mm)", ylab = "")
  mtext(expression(paste("Temperature (", degree, "C)")), side = 2, line = 1.8)
  plot(gamDay, select = 2, se = FALSE, rug = FALSE, las = 1, ylim = c(0, 32), main = "b) Daily data",
       xlab = "Rainfall (mm)",
       yaxt = "n", ylab = "")
  axis(2, labels = NA)
  if(i == 1) dev.off()
```

```
}
```



4.2 Random forest

Daily RF

```
RESPONSE <- "NAccidPer100000"
#predictors
v <- c("HNSaturday", "Holiday", "Month", "Year", "Rain", "Temperature", "Weekday")

DATAanoNA <- na.omit(DtrainDay[,c(RESPONSE, v)])
set.seed(10000)
ran <- RfDay <- ranger(dependent.variable.name = RESPONSE, data = DATAanoNA,
                       importance = 'impurity_corrected',
                       min.node.size = 5, respect.unordered.factors = 'partition',
                       num.trees = 500)

#for predictions
ran2 <- ran2Day <- ranger(dependent.variable.name = RESPONSE, data = DATAanoNA,
                          # importance = 'impurity_corrected',
                          min.node.size = 5, respect.unordered.factors = 'partition',
                          num.trees = 500)

print(ran)

## Ranger result
##
## Call:
## ranger(dependent.variable.name = RESPONSE, data = DATAanoNA, importance = "impurity_corrected", min.node.size = 5, respo
##
## Type: Regression
## Number of trees: 500
## Sample size: 4383
## Number of independent variables: 7
## Mtry: 2
## Target node size: 5
## Variable importance mode: impurity_corrected
## Splitrule: variance
## OOB prediction error (MSE): 0.0417
## R squared (OOB): 0.464

# ranimp <- importance_pvalues(ran, method = "altmann",
#                               num.permutations = 500,
#                               formula = as.formula(paste(RESPONSE, "~")),
#                               data = DATAanoNA)
# ranimp <- ranimp[order(ranimp[,1]),]
```

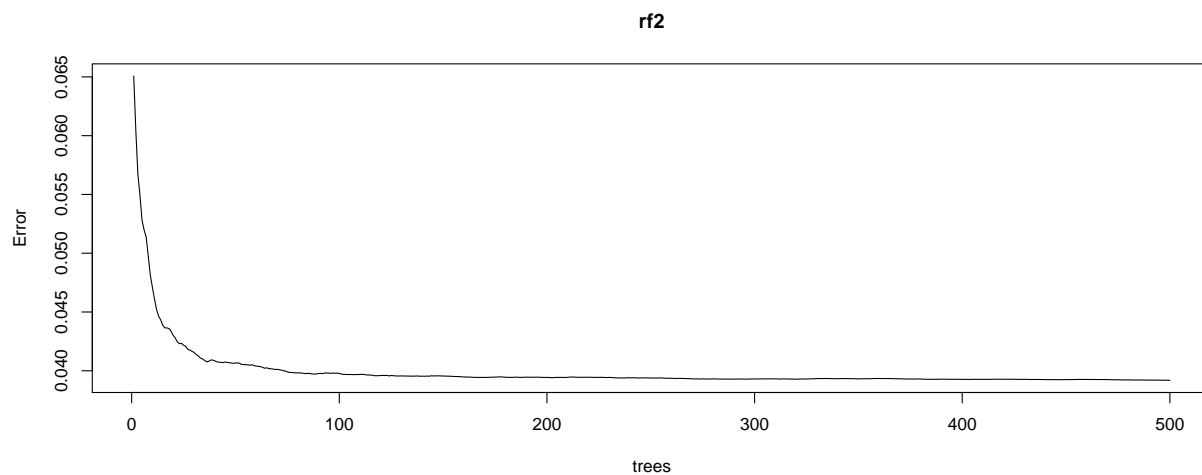
```
# ranimp

set.seed(300000)
rf2 <- rf2Day <- randomForest(y = DATAoNA[,RESPONSE],
                             x = DATAoNA[, v],
                             nodesize = ran$min.node.size,
                             mtry = ran$mtry,
                             ntree = ran$num.trees)

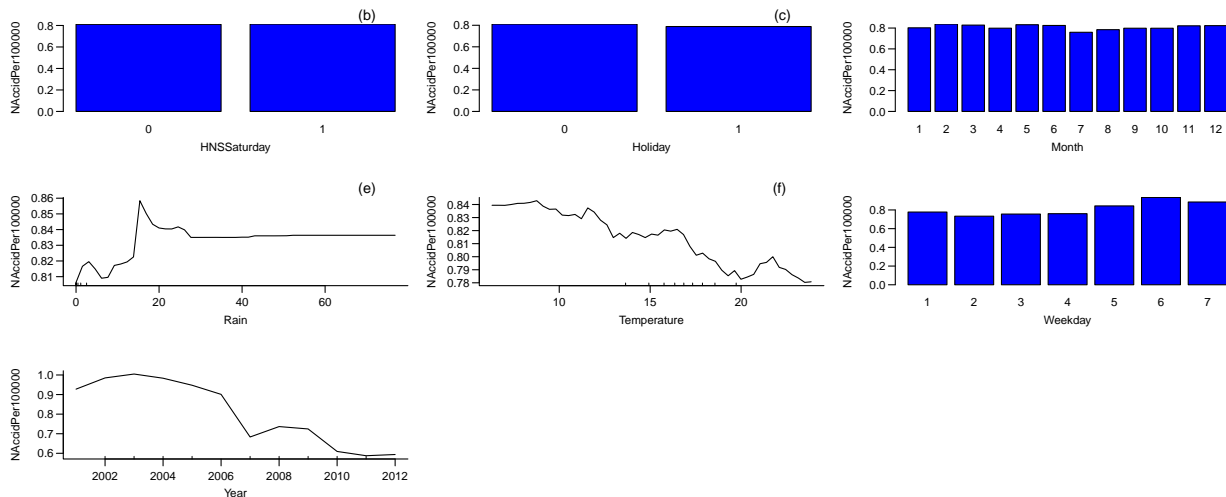
print(rf2)

##
## Call:
## randomForest(x = DATAoNA[, v], y = DATAoNA[, RESPONSE], ntree = ran$num.trees,      mtry = ran$mtry, nodesize = ran$min.no
##           Type of random forest: regression
##           Number of trees: 500
## No. of variables tried at each split: 2
##
##           Mean of squared residuals: 0.0392
##           % Var explained: 49.6

plot(rf2)
```



```
RF <- rf2
preds <- sort(rownames(rf2$importance)) # sort(v)
par(mfrow = c(ceiling(length(preds)/3), 3))
par(bty = "L", mar = c(5, 4, 1, 1) + 0.1, mgp = c(2, 0.7, 0))
for(i in 1:length(preds)) {
  partialPlot(RF, pred.data = DtrainDay, x.var = preds[i],
              las = 1, xlab = preds[i], ylab = "", main = "", xpd = F)
  mtext("NAccidPer100000", side = 2, line = 3, cex = 0.7)
  mtext(paste("(", letters[i], ")", sep = ""), side = 3, line = 0.1, cex = 0.8, adj = -0.37)
}
```



Hourly RF

```
RESPONSE <- "NAccidPer100000"
```

```
#predictors
```

```
v <- c("HNSSaturday", "Holiday", "Hour", "Month", "Year", "Rain", "Temperature",  
      "Rain.l1", "Temperature.l1", "Weekday")
```

```
DATAnoNA <- na.omit(DtrainHour[,c(RESPONSE, v)])
```

```
set.seed(10000)
```

```
ran <- RFHour <- ranger(dependent.variable.name = RESPONSE, data = DATAnoNA,  
                        importance = 'impurity_corrected',  
                        min.node.size = 5, respect.unordered.factors = 'partition',  
                        num.trees = 100)
```

```
## Growing trees.. Progress: 12%. Estimated remaining time: 4 minutes, 9 seconds.  
## Growing trees.. Progress: 26%. Estimated remaining time: 3 minutes, 13 seconds.  
## Growing trees.. Progress: 41%. Estimated remaining time: 2 minutes, 23 seconds.  
## Growing trees.. Progress: 56%. Estimated remaining time: 1 minute, 43 seconds.  
## Growing trees.. Progress: 71%. Estimated remaining time: 1 minute, 6 seconds.  
## Growing trees.. Progress: 84%. Estimated remaining time: 38 seconds.  
## Growing trees.. Progress: 100%. Estimated remaining time: 0 seconds.
```

```
#for predictions
```

```
ran2 <- ran2Hour <- ranger(dependent.variable.name = RESPONSE, data = DATAnoNA,  
                           # importance = 'impurity_corrected',  
                           min.node.size = 5, respect.unordered.factors = 'partition',  
                           num.trees = 100)
```

```
## Growing trees.. Progress: 18%. Estimated remaining time: 2 minutes, 21 seconds.  
## Growing trees.. Progress: 40%. Estimated remaining time: 1 minute, 34 seconds.  
## Growing trees.. Progress: 61%. Estimated remaining time: 1 minute, 0 seconds.  
## Growing trees.. Progress: 83%. Estimated remaining time: 26 seconds.
```

```
print(ran)
```

```
## Ranger result
```

```
##
```

```
## Call:
```

```
##   ranger(dependent.variable.name = RESPONSE, data = DATAnoNA, importance = "impurity_corrected", min.node.size = 5, respo
```

```
##
```

```
## Type: Regression
```

```
## Number of trees: 100
```

```
## Sample size: 105095
```

```
## Number of independent variables: 10
```

```
## Mtry: 3
```

```
## Target node size: 5
```

```
## Variable importance mode: impurity_corrected
```

```
## Splitrule:                variance
## OOB prediction error (MSE): 0.00137
## R squared (OOB):          0.139

# ranimp <- importance_pvalues(ran, method = "altmann",
#                             num.permutations = 500,
#                             formula = as.formula(paste(RESPONSE, ".", sep = " ~ ")),
#                             data = DATAoNA)
# ranimp <- ranimp[order(ranimp[,1]),]
# ranimp

set.seed(300000)
rf2 <- rf2Hour <- randomForest(y = DATAoNA[,RESPONSE],
                              x = DATAoNA[, v],
                              nodesize = ran$min.node.size,
                              mtry = ran$mtry,
                              ntree = ran$num.trees)

print(rf2)

##
## Call:
## randomForest(x = DATAoNA[, v], y = DATAoNA[, RESPONSE], ntree = ran$num.trees,      mtry = ran$mtry, nodesize = ran$min.no
##           Type of random forest: regression
##           Number of trees: 100
## No. of variables tried at each split: 3
##
##           Mean of squared residuals: 0.00137
##           % Var explained: 13.6

plot(rf2)
```




Figure 6: Relative importance of the variables in random forests

```
for(i in 1:2){
  if(i == 1) pdf("./figures/RFimp.pdf", width = 8, height = 3.7)
  par(mar = c(3.3, 5.5, 1.1, 1), mgp = c(2.0, -0.3, 0))
  par(mfrow = c(1, 2))
  #
  tmp <- sort(RFHour$variable.importance)
  names(tmp)[grep("HNS", names(tmp))] <- "HNCS"
  names(tmp)[grep("ture.", names(tmp))] <- "Temperature(t-1)"
  names(tmp)[grep("ain.", names(tmp))] <- "Rainfall(t-1)"
  names(tmp)[names(tmp) == "Rain"] <- "Rainfall"
  barplot(tmp,
    beside = TRUE, las = 1, xlim = c(-2, 10),
    main = "a) Hourly data",
    xlab = "Importance",
```

```

        col = COL[4], border = NA, cex.names = 0.8, xaxt = "n",
        horiz = TRUE)
par(mgp = c(2.0, 0.8, 0))
axis(1)
#
tmp <- sort(RFDay$variable.importance)
names(tmp)[grep("HNS", names(tmp))] <- "HNCS"
names(tmp)[names(tmp) == "Rain"] <- "Rainfall"
barplot(tmp,
        beside = TRUE, las = 1, xlim = c(0, 100),
        main = "b) Daily data",
        xlab = "Importance",
        col = COL[4], border = NA, cex.names = 0.8,
        horiz = TRUE)
if(i == 1) dev.off()
}

```

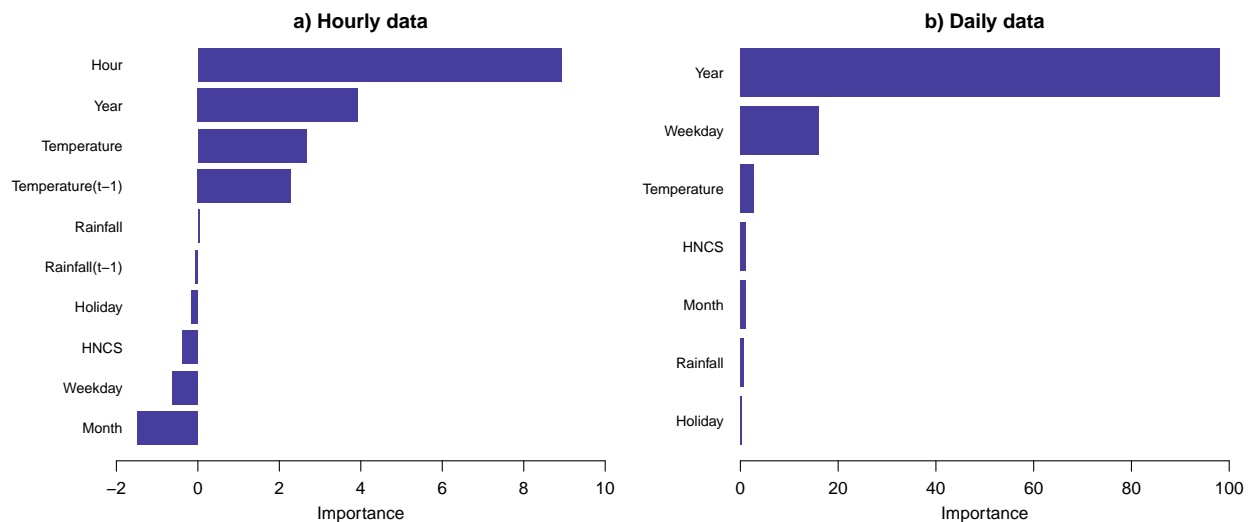


Figure 7: Partial dependence interaction plots from the random forests

```

pdf("./figures/RFinter.pdf", width = 5, height = 4.7)
par(mfrow = c(1, 2))
par(mar = c(3.3, 3, 1.1, 0.1), mgp = c(2.0, 0.8, 0))
v <- c("HNSSaturday", "Holiday", "Hour", "Month", "Year", "Rain", "Temperature",
       "Rain.l1", "Temperature.l1", "Weekday")
DATAanoNA <- na.omit(DtrainHour[,c(RESPONSE, v)])
RF <- rf2Hour
plotmo(RF, pmethod = "partdep", ylim = c(0, 32),
       all1 = FALSE,
       degree1 = FALSE,
       degree2 = c("Rain", "Temperature"),
       type2 = "contour",
       caption = "", main = "a) Hourly data",
       all2 = TRUE)

## calculating partdep for Rain:Temperature 01234567890

v <- c("HNSSaturday", "Holiday", "Month", "Year", "Rain", "Temperature", "Weekday")
DATAanoNA <- na.omit(DtrainDay[,c(RESPONSE, v)])
RF <- rf2Day
plotmo(RF, pmethod = "partdep", ylim = c(0, 32),
       all1 = FALSE,
       degree1 = FALSE,
       degree2 = c("Rain", "Temperature"),
       type2 = "contour",
       caption = "", main = "b) Daily data",
       all2 = TRUE)

```



```
## calculating partdep for Rain:Temperature 01234567890
dev.off()
## pdf
## 2
```

4.3 Performance evaluation

Number of accidents per year in the training set

```
sum(DtrainDay$NAccidPer100000) / length(unique(DtrainDay$Year))
## [1] 296
```

Test daily

```
Dtest <- DtestDay
gamfit <- gamDay
ran2 <- ran2Day

pred_gam <- predict(gamfit, Dtest)
pred_ran <- predict(ran2, Dtest)$predictions
#PMAE
mean(abs(Dtest[,RESPONSE] - pred_gam))
## [1] 0.189

mean(abs(Dtest[,RESPONSE] - pred_ran))
## [1] 0.113

#PRMSE
sqrt(mean((Dtest[,RESPONSE] - pred_gam)^2))
## [1] 0.223

sqrt(mean((Dtest[,RESPONSE] - pred_ran)^2))
## [1] 0.143

#PMAPE
100*mean(abs((Dtest[,RESPONSE] - pred_gam)/Dtest[,RESPONSE]))
## [1] 30.4

100*mean(abs((Dtest[,RESPONSE] - pred_ran)/Dtest[,RESPONSE]))
## [1] 21.2

tapply(Dtest[,RESPONSE], Dtest$Year, sum)
## 2013 2014 2015
## 235 215 179

tapply(pred_gam, Dtest$Year, sum)
## 2013 2014 2015
## 174 152 116

tapply(pred_ran, Dtest$Year, sum)
## 2013 2014 2015
## 216 217 198
```

Test hourly

```
Dtest <- DtestHour
gamfit <- gamHour
```

```

ran2 <- ran2Hour

pred_gam <- predict(gamfit, Dtest)
pred_ran <- predict(ran2, Dtest)$predictions
#PMAE
mean(abs(Dtest[,RESPONSE] - pred_gam))

## [1] 0.0188

mean(abs(Dtest[,RESPONSE] - pred_ran))

## [1] 0.0191

#PRMSE
sqrt(mean((Dtest[,RESPONSE] - pred_gam)^2))

## [1] 0.0257

sqrt(mean((Dtest[,RESPONSE] - pred_ran)^2))

## [1] 0.0246

#PMAPE
100*mean(abs((Dtest[,RESPONSE] - pred_gam)/Dtest[,RESPONSE]))

## [1] Inf

100*mean(abs((Dtest[,RESPONSE] - pred_ran)/Dtest[,RESPONSE]))

## [1] Inf

tapply(Dtest[,RESPONSE], Dtest$Year, sum)

## 2013 2014 2015
## 235 215 179

tapply(pred_gam, Dtest$Year, sum)

## 2013 2014 2015
## 176 155 121

tapply(pred_ran, Dtest$Year, sum)

## 2013 2014 2015
## 218 216 197

```

Average annual accident rate in the training period

```

tmp = tapply(DtrainDay[,RESPONSE], DtrainDay$Year, sum)
tmp

## 2001 2002 2003 2004 2005 2006 2007 2008 2009 2010 2011 2012
## 335 371 385 371 345 326 240 281 267 216 200 212

mean(tmp)

## [1] 296

```

Save all objects from the R environment:

```

save.image(file = "./dataderived/image_MexAnalysis.RData")

```